


Texture Mapping



Oregon State
University

Mike Bailey

mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nc-nd/4.0/)



Oregon State
University


Computer Graphics

Also, to prevent confusion, the texture pixels are not called *pixels*. A pixel is a dot in the final screen image. A dot in the texture image is called a *texture element*, or *texel*.

Similarly, to avoid terminology confusion, a texture's width and height dimensions are not called X and Y . They are called **S** and **T**.


A texture map is not generally indexed by its actual resolution coordinates. Instead, it is indexed by a coordinate system that is resolution-independent. The left side is always **S=0.**, the right side is **S=1.**, the bottom is **T=0.**, and the top is **T=1.**

Thus, you do not need to be aware of the texture's resolution when you are specifying coordinates that point into it. Think of S and T as a measure of what fraction of the way you are into the texture.




Texture mapping is a computer graphics operation in which a separate image, referred to as the **texture**, is stretched onto a piece of 3D geometry and follows it however it is transformed. This image is also known as a **texture map**.

This can be most any image. Some graphics hardware requires the image's pixel dimensions to be a **power of two**. (This restriction has been lifted on most graphics cards, but just to be safe....) The X and Y dimensions do not need to be the *same* power of two, just a power of two. So, a 128x512 image would be OK, a 129x511 image might not.



The mapping between the geometry of the **3D object** and the S and T of the **texture image** works like this:



You specify an (s,t) pair at each vertex, along with the vertex coordinate. At the same time that OpenGL is interpolating the coordinates, colors, etc. inside the polygon, it is also interpolating the (s,t) coordinates. Then, when OpenGL goes to draw each pixel, it uses that pixel's interpolated (s,t) to lookup a color in the texture image.

Enable texture mapping:

```
glEnable( GL_TEXTURE_2D );
```

Draw your polygons, specifying **s** and **t** at each vertex:

```
glBegin( GL_POLYGON );  
    glTexCoord2f( s0, t0 );  
    glNormal3f( nx0, ny0, nz0 );  
    glVertex3f( x0, y0, z0 );  
  
    glTexCoord2f( s1, t1 );  
    glNormal3f( nx1, ny1, nz1 );  
    glVertex3f( x1, y1, z1 );  
  
    ...  
glEnd( );
```

If this geometry is static (i.e., will never change), *it is a good idea to put this all into a display list.*


Disable texture mapping:

```
glDisable( GL_TEXTURE_2D );
```



```
glTexCoord2f( s0, t0 );
```

The easiest way to figure out what s and t are at a particular vertex is to figure out what fraction across the object the vertex is living at. For a plane, this is pretty easy:




$$s = \frac{x - Xmin}{Xmax - Xmin} \quad t = \frac{y - Ymin}{Ymax - Ymin}$$



`glTexCoord2f(s0, t0);`

Or, for a sphere, you do the same thing you did for the plane, only the interpolated variables are angular (spherical) coordinates instead of linear coordinates

$$s = \frac{\Theta - (-\pi)}{2\pi} \quad t = \frac{\Phi - (-\frac{\pi}{2})}{\pi}$$



The Sphere code does it like this:


`s = (lng + M_PI) / (2.*M_PI);`

`t = (lat + M_PI/2.) / M_PI;`




```
glTexCoord2f( s0, t0 );
```

Uh-oh. Now what? Here's where it gets tougher...,



$s = ?$

$t = ?$






You really are at the mercy of whoever did the modeling and assigned the s,t coordinates... 9



Oregon State
University

Computer Graphics



```
unsigned char *BmpToTexture( char *, int *, int * );  
  
unsigned char *Texture;  
int width, height;  
...  
Texture = BmpToTexture( "filename.bmp", &width, &height );
```

This function is found in your sample code. The BMP file *filename.bmp* needs to be created by something that writes ***uncompressed 24-bit color BMP files***, or converted to the uncompressed BMP format by a tool such as ImageMagick's *convert*, Adobe *Photoshop*, or *GIMP*.

Note: this function should be called *once*, and called from `InitGraphics()`. Do not call it in the `Display()` function.




Define the texture wrapping parameters. This will control what happens when a texture coordinate is greater than 1.0 or less than 0.0:

```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, wrap );  
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, wrap );
```

where *wrap* is:

GL_REPEAT specifies that this pattern will repeat (i.e., wrap-around) if transformed texture coordinates less than 0.0 or greater than 1.0 are encountered.

GL_CLAMP specifies that the pattern will “stick” to the value at 0.0 or 1.0.



Define the texture filter parameters. This will control what happens when a texture is scaled up or down.


```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, filter );  
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, filter );
```

where *filter* is:


GL_NEAREST specifies that point sampling is to be used when the texture map needs to be magnified or minified.

GL_LINEAR specifies that bilinear interpolation among the four nearest neighbors is to be used when the texture map needs to be magnified or minified.

GL_NEAREST



GL_LINEAR



This tells OpenGL what to do with the texel colors when it gets them:


```
glTexEnvf( GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE, mode );
```

There are several *modes* that can be used. Two of the most useful are:


GL_REPLACE specifies that the 3-component texture will be applied as an opaque image on top of the polygon, replacing the polygon's specified color.

GL_MODULATE specifies that the 3-component texture will be applied as piece of colored plastic on top of the polygon. The polygon's specified color "shines" through the plastic texture. This is very useful for applying lighting to textures: paint the polygon white with lighting and let it shine up through a texture.

GL_REPLACE



GL_MODULATE



```
int width, height;
...
unsigned char *Texture = BmpToTexture( "filename.bmp", &width, &height );
int level, ncomps, border;
...
glPixelStorei( GL_UNPACK_ALIGNMENT, 1 );
glTexImage2D( GL_TEXTURE_2D, level, ncomps, width, height, border, GL_RGB, GL_UNSIGNED_BYTE, Texture );
```

where:

level	is used with mip-mapping. Use 0 for now.
ncomps	number of components in this texture: 3 if using RGB, 4 if using RGBA.
width	width of this texture map, in pixels.
height	height of this texture map, in pixels.
border	width of the texture border, in pixels. Use 0 for now.
Texture	the name of an array of unsigned characters holding the texel colors.



Oreg

University

Computer Graphics

This function physically **transfers** the array of texels from the CPU to the GPU and makes it the current active texture. You can get away with specifying this ahead of time only if you are using a **single texture**. If you are using multiple textures, you must make each current in **Display()** right before you need it. See the section below about *binding* textures.

In addition to the Projection and ModelView matrices, OpenGL maintains a transformation for texture map coordinates **S** and **T** as well. You use all the same transformation routines you are used to: **glRotatef()**, **glScalef()**, **glTranslatef()**, but you must first specify the **Matrix Mode**:

```
glMatrixMode( GL_TEXTURE );
```

The only trick to this is to remember that you are transforming the *texture coordinates*, not the *texture image*. Transforming the texture image forward is the same as transforming the texture coordinates backwards:


Scale = 0.5



Scale = 1.




Scale = 2.




The only trick to this is to remember that you are transforming the texture coordinates, not the texture image. Transforming the texture image forward is the same as transforming the texture coordinates backwards:

Angle = -45.



Angle = 0.



Angle = 45.



The OpenGL **glTexImage2D** function doesn't just use that texture, it **downloads** it from the CPU to the GPU, *every time that call is made!* After the download, this texture becomes the “current texture image”.

```
glTexImage2D( GL_TEXTURE_2D, level, ncomps, width, height, border, GL_RGB, GL_UNSIGNED_BYTE, Texture );
```

If your scene has only one texture, this is easy to manage. Just do it once and forget about it.

But, if you have several textures, all to be used at different times on different objects, it will be important to maximize the efficiency of how you create, store, and manage those textures. In this case you should bind **texture objects**.

Texture objects leave your textures on the graphics card and then re-uses them, which is always going to be faster than re-loading them.



Create a texture object by generating a texture name and then bind the texture object to the texture data and texture properties. The first time you execute **glBindTexture()**, you fill the texture object. Subsequent times you do this, you are making that texture object current. So, create *global* Texture IDs like this:

```
GLuint Tex0, Tex1;           // global variables
```

```
...
```

Then, at the end of **InitGraphics()** you add:

```
int width, height;
```

```
unsigned char * TextureArray0 = BmpToTexture( "file.bmp", &width, &height );
```

```
glPixelStorei( GL_UNPACK_ALIGNMENT, 1 );
```

```
glGenTextures( 1, &Tex0 );           // assign binding "handles"
```

```
glGenTextures( 1, &Tex1 );
```

```
...
```

```
glBindTexture( GL_TEXTURE_2D, Tex0 ); // make the Tex0 texture current
```

```
// and set its parameters
```

```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_CLAMP );
```

```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_CLAMP );
```

```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_LINEAR );
```

```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_LINEAR );
```

```
glTexEnvf( GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE, GL_REPLACE );
```

```
glTexImage2D( GL_TEXTURE_2D, 0, 3, width, height, 0, GL_RGB, GL_UNSIGNED_BYTE, TextureArray0 );
```

```
glBindTexture( GL_TEXTURE_2D, tex1 );           // designate the tex1 texture as the “current texture”  
                                                // and set its parameters
```

```
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_REPEAT );  
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_REPEAT );  
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_NEAREST );  
glTexParameteri( GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_NEAREST );  
glTexEnvf( GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE, GL_MODULATE );  
glTexImage2D( GL_TEXTURE_2D, 0, 3, 512, 512, 0, GL_RGB, GL_UNSIGNED_BYTE, TextureArray1 );
```

Then, later on in Display():

```
glEnable( GL_TEXTURE_2D );
```

```
glBindTexture( GL_TEXTURE_2D, tex0 );  
glBegin( GL_QUADS );
```

...

```
glEnd( );
```

```
glBindTexture( GL_TEXTURE_2D, tex1 );  
glBegin( GL_TRIANGLE_STRIP );
```

...

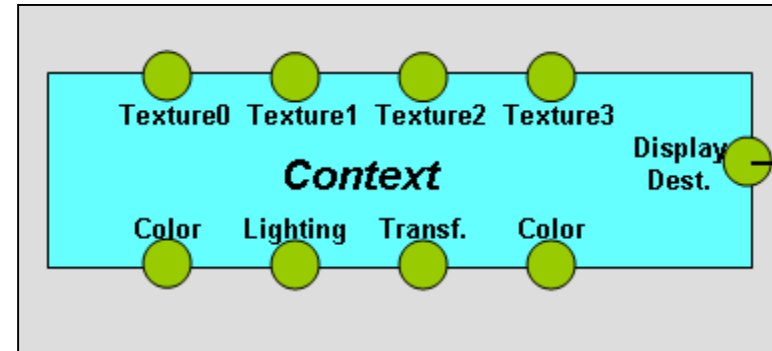
```
glEnd( );
```



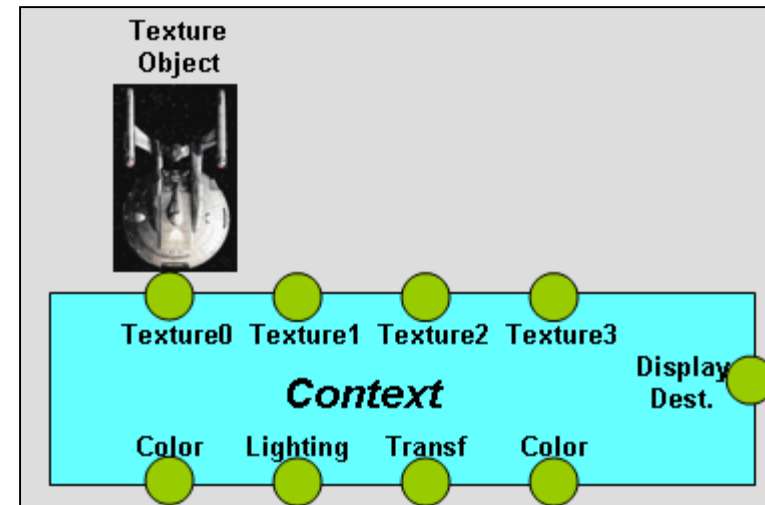
The OpenGL Rendering Context contains all the characteristic information necessary to produce an image from geometry. This includes the current transformations, colors, lighting, textures, where to send the display, etc.

The OpenGL term “binding” refers to “attaching” or “docking” (a metaphor which I find to be more visually pleasing) an OpenGL object to the Context. You can then assign characteristics, and they will “flow” through the Context into the object.

Before
Binding



After
Binding



Some Great Uses for Texture Mapping you have seen in the Movies



Disney



Disney



Disney

Yes, I know, I know, these are older examples, but I especially like them because, at the time, the CG (and the textures) became part of the story-telling.



Some Great Uses for Texture Mapping you have seen in the Movies



Disney



Disney



Pixar



Disney




Oregon State
University

Computer Graphics

Bonus Topic: Procedural Texture Mapping

You can also create a texture from data on-the-fly. In this case, the fragment shader takes a grid of heights and uses cross-products to produce surface normal vectors for lighting.

While this is “procedural”, the amount of height data is finite, so you can still run out of resolution



We cover this more in the
shaders course: CS 457/557

Although this looks like an incredible amount of polygonal scene detail, the geometry for this scene consists of just a ***single quadrilateral***



Bonus Topic: Procedural Texture Mapping

“Mandelzoom”:

In this case, the texture is a pure equation, so you never run out of resolution. (You do run out of floating-point precision, however.)

We cover this more in the
shaders course: CS 457/557

